



deeplearning.ai

# Probability and Bayes' Rule

# Outline

- Probabilities
- Bayes' rule (Applied in different fields, including NLP)
- Build your own Naive-Bayes tweet classifier!

# Introduction

Corpus of tweets

		Positive		

Tweets containing the word  
“happy”

		Positive		

# Probabilities

Corpus of tweets

		Positive		
		Negative		

$A \rightarrow \text{Positive tweet}$

$$P(A) = P(\text{Positive}) = N_{\text{pos}} / N$$

# Probabilities

Corpus of tweets

		Positive		
		Negative		

$A \rightarrow \text{Positive tweet}$

$$P(A) = N_{\text{pos}} / N = 13 / 20 = 0.65$$

$$P(\text{Negative}) = 1 - P(\text{Positive}) = 0.35$$

# Probabilities

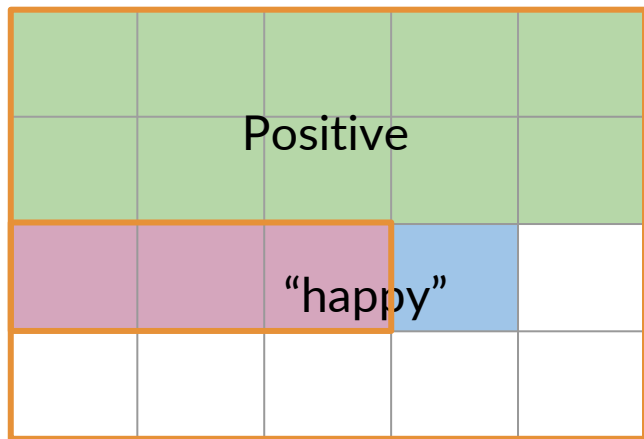
Tweets containing the word  
"happy"


$B \rightarrow$  tweet contains "happy".

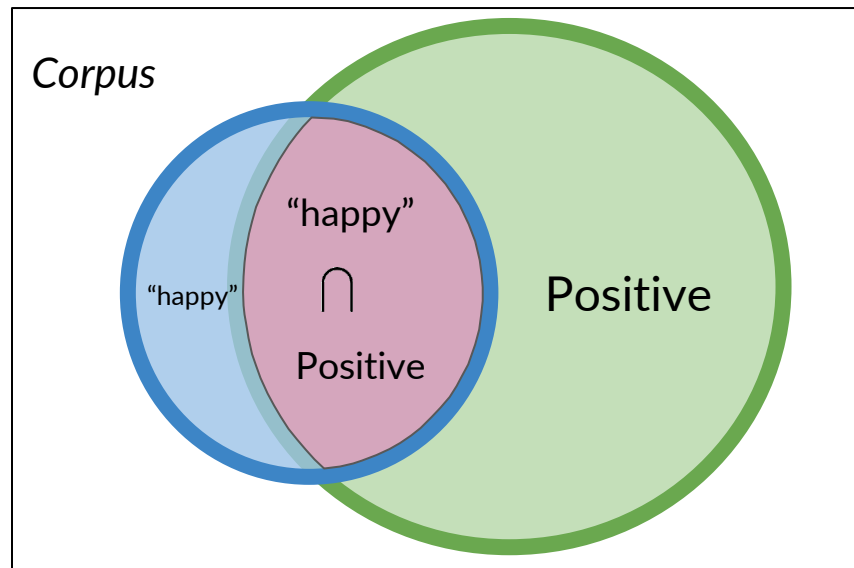
$$P(B) = P(\text{happy}) = N_{\text{happy}} / N$$

$$P(B) = 4 / 20 = 0.2$$

# Probability of the intersection



$$P(A \cap B) = P(A, B) = \frac{3}{20} = 0.15$$



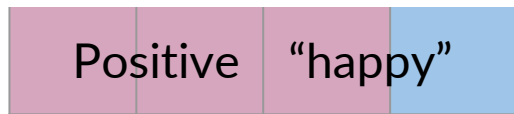


deeplearning.ai

# Bayes' Rule

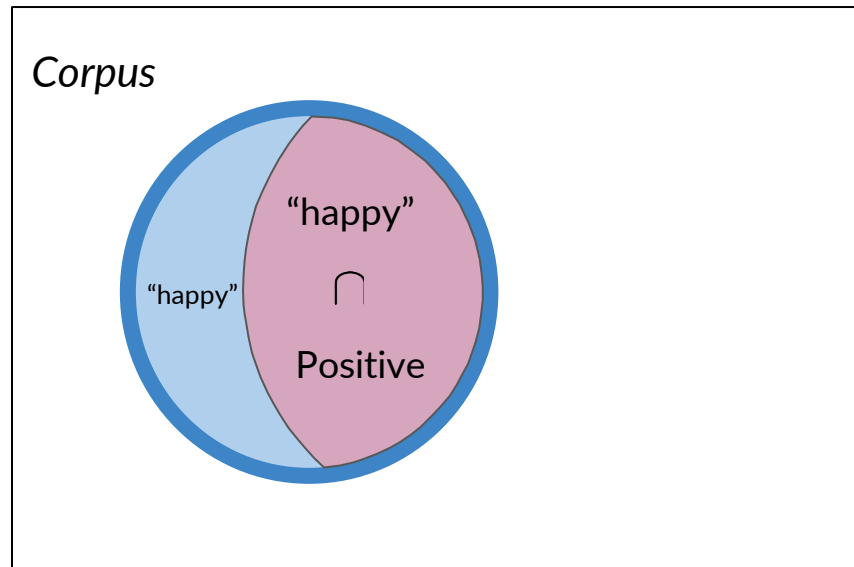


# Conditional Probabilities

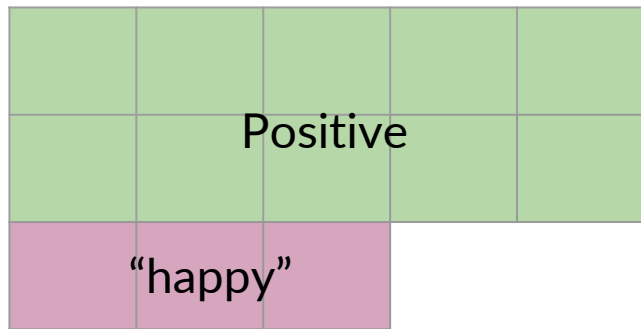


$$P(A | B) = P(\text{Positive} | \text{"happy"})$$

$$P(A | B) = 3 / 4 = 0.75$$

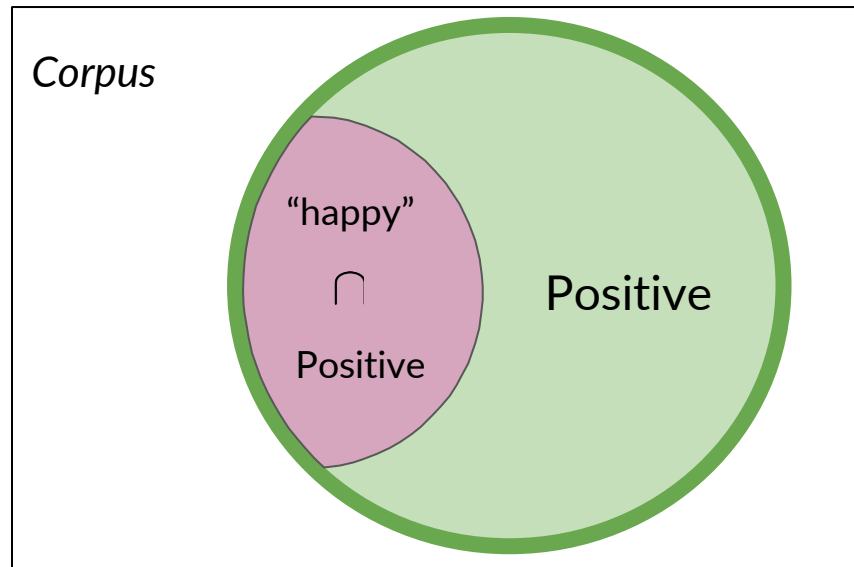


# Conditional Probabilities



$$P(B | A) = P(\text{“happy”} | \text{Positive})$$

$$P(B | A) = 3 / 13 = 0.231$$



# Conditional probabilities

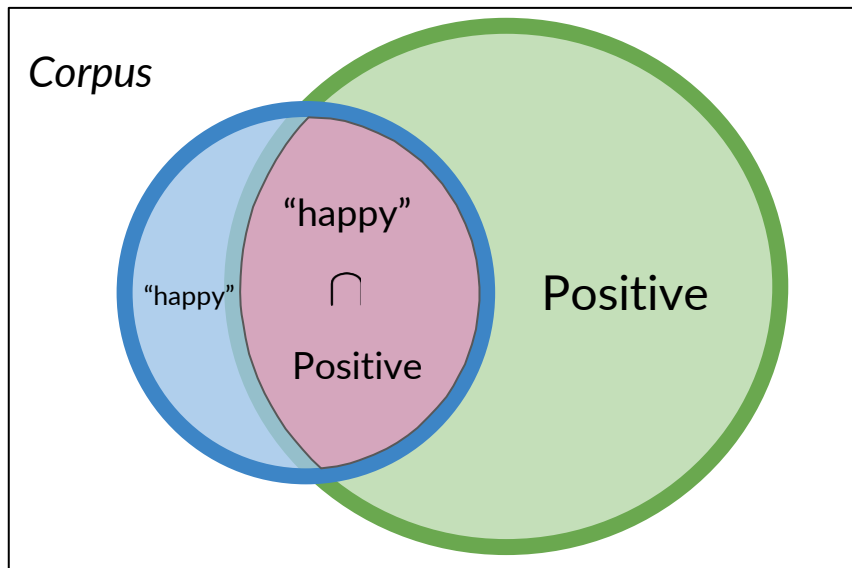


```
graph TD; A[Conditional probabilities] --> B[Probability of B, given A happened]; A --> C[Looking at the elements of set A, the chance that one also belongs to set B]
```

Probability of B, given A happened

Looking at the elements of set A,  
the chance that one also belongs to  
set B

# Conditional probabilities



$$P(\text{Positive} | \text{"happy"}) = \frac{P(\text{Positive} \cap \text{"happy"})}{P(\text{"happy"})}$$

## Bayes' rule

$$P(\text{Positive} | \text{"happy"}) = \frac{P(\text{Positive} \cap \text{"happy"})}{P(\text{"happy"})}$$

$$P(\text{"happy"} | \text{Positive}) = \frac{P(\text{"happy"} \cap \text{Positive})}{P(\text{Positive})}$$

## Quiz

**Objective:** Derive Bayes' rule from the equations given on the last slide.

**Question:**

From the equations presented below, express the probability of a tweet being positive given that it contains the word happy in terms of the probability of a tweet containing the word happy given that it is positive

$$P(\text{Positive} | \text{"happy"}) = \frac{P(\text{Positive} \cap \text{"happy"})}{P(\text{"happy"})} \quad P(\text{"happy"} | \text{Positive}) = \frac{P(\text{"happy"} \cap \text{Positive})}{P(\text{Positive})}$$

**Type:** Multiple Choice, single answer

**Options and solution:**

$$P(\text{Positive} | \text{"happy"}) = P(\text{"happy"} | \text{Positive}) \times \frac{P(\text{Positive})}{P(\text{"happy"})}$$

That's right. You just derived Bayes' rule.

$$P(\text{Positive} | \text{"happy"}) = P(\text{"happy"} | \text{Positive}) \times \frac{P(\text{"happy"})}{P(\text{Positive})}$$

The ratio is upside-down in this equation.

$$P(\text{Positive} | \text{"happy"}) = P(\text{"happy"} \cap \text{Positive}) \times \frac{P(\text{Positive})}{P(\text{"happy"})}$$

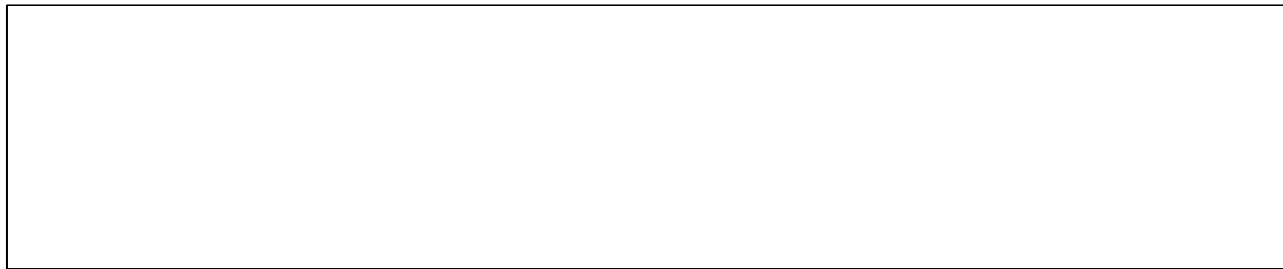
Your result should not include any intersection probabilities.

$$P(\text{Positive} | \text{"happy"}) = P(\text{"happy"} \cap \text{Positive}) \times \frac{P(\text{"happy"})}{P(\text{Positive})}$$

Your result should not include any intersection probabilities.

## Bayes' rule

$$P(\text{Positive} | \text{"happy"}) = P(\text{"happy"} | \text{Positive}) \times \frac{P(\text{Positive})}{P(\text{"happy"})}$$



## Quiz: Bayes' Rule Applied

**Objective:** Compute conditional probability using Bayes Rule

**Question:**

$$P(X|Y) = P(Y|X) \times \frac{P(X)}{P(Y)}$$

Here, again, is Bayes' rule:

Suppose that in your dataset, 25% of the positive tweets contain the word 'happy'. You also know that a total of 13% of the tweets in your dataset contain the word 'happy', and that 40% of the total number of tweets are positive. You observe the tweet: "happy to learn NLP". What is the probability that this tweet is positive?

**Type:** Multiple Choice, single answer

**Options and solution:**

**A:  $P(\text{Positive} | \text{"happy"}) = 0.77$**  That's right. You just applied Bayes' rule.

**B:  $P(\text{Positive} | \text{"happy"}) = 0.08$**  Oops, looks like you might have the ratio of  $P(X)$  and  $P(Y)$  upside-down.

**C:  $P(\text{Positive} | \text{"happy"}) = 0.10$**  Remember to calculate the ratio in the formula for Bayes' rule.

**D:  $P(\text{Positive} | \text{"happy"}) = 1.92$**  Did you use the probability of a tweet being positive? Remember that a fractional probability must be between 0 and 1.



# Summary

- Conditional probabilities  $\longrightarrow$  Bayes' Rule
- $P(X|Y) = P(Y|X) \times \frac{P(X)}{P(Y)}$



deeplearning.ai

# Naïve Bayes Introduction

---

# Naïve Bayes for Sentiment Analysis

## Positive tweets

I am happy because I am learning NLP  
I am happy, not sad.

## Negative tweets

I am sad, I am not learning NLP  
I am sad, not happy

word	Pos	Neg
I	3	3
am	3	3
happy	2	1
because	1	0
learning	1	1
NLP	1	1
sad	1	2
not	1	2
$N_{\text{class}}$	13	12



$P(w_i \mid \text{class})$

word	Pos	Neg
I	0.24	0.25
am	0.24	0.25
happy	0.15	0.08
because	0.08	0
learning	0.08	0.08
NLP	0.08	0.08
sad	0.08	0.08
not	0.08	0.17

# Naïve Bayes

Tweet: I am happy today; I am learning.

$$\prod_{i=1}^m \frac{P(w_i|pos)}{P(w_i|neg)} = \frac{0.14}{0.10} = 1.4 > 1$$

$$\frac{\cancel{0.20}}{\cancel{0.20}} * \frac{\cancel{0.20}}{\cancel{0.20}} * \frac{0.14}{0.10} * \frac{\cancel{0.20}}{\cancel{0.20}} * \frac{\cancel{0.20}}{\cancel{0.20}} * \frac{\cancel{0.10}}{\cancel{0.10}}$$

word	Pos	Neg
I	0.20	0.20
am	0.20	0.20
happy	0.14	0.10
because	0.10	0.05
learning	0.10	0.10
NLP	0.10	0.10
sad	0.10	0.10
not	0.10	0.15

# Summary

- Naive Bayes inference condition rule for binary classification
- Table of probabilities

$$\prod_{i=1}^m \frac{P(w_i|pos)}{P(w_i|neg)}$$



deeplearning.ai

# Laplacian Smoothing



# Laplacian Smoothing

$$P(w_i | \text{class}) = \frac{\text{freq}(w_i, \text{class})}{N_{\text{class}}} \quad \text{class} \in \{\text{Positive}, \text{Negative}\}$$

$$P(w_i | \text{class}) = \frac{\text{freq}(w_i, \text{class}) + 1}{N_{\text{class}} + V_{\text{class}}}$$

$N_{\text{class}}$  = frequency of all words in class

$V_{\text{class}}$  = number of unique words in class



# Introducing $P(w_i | \text{class})$ with smoothing

word	Pos	Neg		word	Pos	Neg
I	3	3	<div> <math display="block">P(w_i   \text{class}) = \frac{\text{count} + 1}{\text{sum} + V}</math> <p> <math>\frac{3+1}{13+8}</math> </p> </div>			
am	3	3				
happy	2	1				
because	1	0				
learning	1	1				
NLP	1	1				
sad	1	2				
not	1	2	$V = 8$			
Nclass	13			Sum	1	1
12						

# Summary

- Laplacian smoothing to avoid  $P(w_i|class) = 0$
- Naïve Bayes formula

$$\prod_{i=1}^m \frac{P(w_i|pos)}{P(w_i|neg)}$$



deeplearning.ai

# Log Likelihood, Part 1

---

# Ratio of probabilities

Positive  $\infty$

Neutral 1

Negative 0

word	Pos	Neg	ratio
I	0.19	0.20	
am	0.19	0.20	
happy	0.14	0.10	
because	0.10	0.05	
learning	0.10	0.10	
NLP	0.10	0.10	1
sad	0.10	0.15	
not	0.10	0.15	0.6

$$\text{ratio}(w_i) = \frac{P(w_i \mid \text{Pos})}{P(w_i \mid \text{Neg})}$$

$$\approx \frac{\text{freq}(w_i, 1) + 1}{\text{freq}(w_i, 0) + 1}$$

# Naïve Bayes' inference

$class \in \{pos, neg\}$

$w \rightarrow$  Set of  $m$  words in a tweet

$$\frac{P(pos)}{P(neg)} \prod_{i=1}^m \frac{P(w_i|pos)}{P(w_i|neg)} > 1$$

- A simple, fast, and powerful baseline
- A probabilistic model used for classification

# Log Likelihood

$$\frac{P(pos)}{P(neg)} \prod_{i=1}^m \frac{P(w_i|pos)}{P(w_i|neg)}$$

- Products bring risk of underflow

- $\log(a * b) = \log(a) + \log(b)$

- $\log\left(\frac{P(pos)}{P(neg)} \prod_{i=1}^n \frac{P(w_i|pos)}{P(w_i|neg)}\right) \Rightarrow \log \frac{P(pos)}{P(neg)} + \sum_{i=1}^n \log \frac{P(w_i|pos)}{P(w_i|neg)}$

**log prior + log likelihood**



# Calculating Lambda

tweet: I am happy because I am learning.

$$\lambda(w) = \log \frac{P(w|pos)}{P(w|neg)}$$

$$\lambda(am) = \log \frac{0.04}{0.04} = \log(1) = 0$$

word	Pos	Neg	$\lambda$
I	0.05	0.05	
am	0.04	0.04	
happy	0.09	0.01	
because	0.01	0.01	
learning	0.03	0.01	
NLP	0.02	0.02	
sad	0.01	0.09	
not	0.02	0.03	-0.4

# Summing the Lambdas

doc: I am happy because I am learning.

$$\lambda(w) = \log \frac{P(w|pos)}{P(w|neg)}$$

$$\lambda(\text{happy}) = \log \frac{0.09}{0.01} \approx 2.2$$

word	Pos	Neg	$\lambda$
I	0.05	0.05	0
am	0.04	0.04	0
happy	0.09	0.01	
because	0.01	0.01	
learning	0.03	0.01	
NLP	0.02	0.02	
sad	0.01	0.09	
not	0.02	0.03	-0.4

# Summary

- Word sentiment

$$ratio(w) = \frac{P(w|pos)}{P(w|neg)}$$

$$\lambda(w) = \log \frac{P(w|pos)}{P(w|neg)}$$



deeplearning.ai

# Log Likelihood, Part 2

# Log Likelihood

doc: I am happy because I am learning.

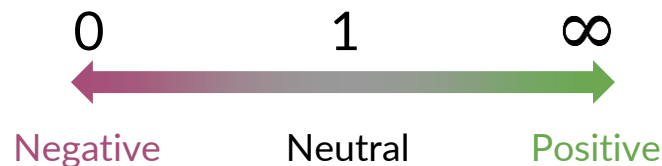
$$\sum_{i=1}^m \log \frac{P(w_i|pos)}{P(w_i|neg)} = \sum_{i=1}^m \lambda(w_i)$$

$$\text{log likelihood} = 0 \begin{matrix} + \\ 0 \end{matrix} \begin{matrix} + \\ 2.2 \end{matrix} \begin{matrix} + \\ 0 \end{matrix} \begin{matrix} + \\ 0 \end{matrix} \begin{matrix} + \\ 0 \end{matrix} \begin{matrix} + \\ 1.1 \end{matrix} = 3.3$$

word	Pos	Neg	$\lambda$
I	0.05	0.05	0
am	0.04	0.04	0
happy	0.09	0.01	2.2
because	0.01	0.01	0
learning	0.03	0.01	1.1
NLP	0.02	0.02	0
sad	0.01	0.09	-2.2
not	0.02	0.03	-0.4

# Log Likelihood

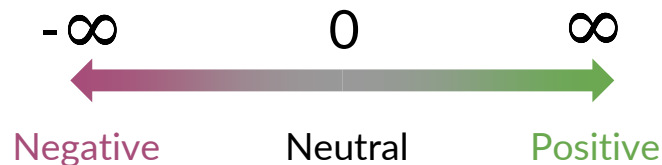
$$\prod_{i=1}^m \frac{P(w_i|pos)}{P(w_i|neg)} > 1$$



$$3.3 > 0$$



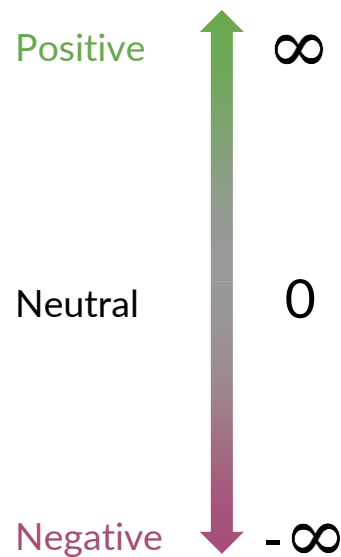
$$\sum_{i=1}^m \log \frac{P(w_i|pos)}{P(w_i|neg)} > 0$$



# Summary

Tweet sentiment:

$$\log \prod_{i=1}^m \text{ratio}(w_i) = \sum_{i=1}^m \lambda(w_i) > 0$$





deeplearning.ai

# Training Naïve Bayes



# Outline

- Five steps for training a Naïve Bayes model

# Training Naïve Bayes

Step 0: Collect and annotate corpus

Positive tweets

I am happy because I am learning

NLP  
I am happy, not sad. @NLP

Negative tweets

I am sad, I am not learning NLP

I am sad, not happy!!

Step 1:  
Preprocess

- Lowercase
- Remove punctuation, urls, names
- Remove stop words
- Stemming
- Tokenize sentences

Positive tweets

[happi, because, learn, NLP]

[happi, not, sad]

Negative tweets

[sad, not, learn, NLP]

[sad, not, happi]

# Training Naïve Bayes

## Positive tweets

[happi, because, learn, NLP]

[happi, not, sad]

## Negative tweets

[sad, not, learn, NLP]

[sad, not, happi]

Step 2:  
Word  
count

freq(w, class)

word	Pos	Neg
happi	2	1
because	1	0
learn	1	1
NLP	1	1
sad	1	2
not	1	2
$N_{\text{class}}$	7	7

# Training Naïve Bayes

freq(w, class)		
word	Pos	Neg
happi	2	1
because	1	0
learn	1	1
NLP	1	1
sad	1	2
not	1	2
$N_{\text{class}}$	7	7

Step 3:  
 $P(w|\text{class})$

$$V_{\text{class}} = 6$$

$$\frac{\text{freq}(w, \text{class}) + 1}{N_{\text{class}} + V_{\text{class}}}$$

$$\lambda(w) = \log \frac{P(w|\text{pos})}{P(w|\text{neg})}$$

Step 4:  
Get  
lambda

word	Pos	Neg	$\lambda$
happy	0.23	0.15	0.43
because	0.15	0.07	0.6
learning	0.08	0.08	0
NLP	0.08	0.08	0
sad	0.08	0.17	-0.75
not	0.08	0.17	-0.75

# Training Naïve Bayes

Step 5:  
Get the  
log prior

$D_{\text{pos}}$  = Number of positive tweets  
 $D_{\text{neg}}$  = Number of negative tweets

$$\text{logprior} = \log \frac{D_{\text{pos}}}{D_{\text{neg}}}$$

If dataset is balanced,  $D_{\text{pos}} = D_{\text{neg}}$  and  $\text{logprior} = 0$ .

# Summary

1. Get or annotate a dataset with positive and negative tweets
2. Preprocess the tweets:  $\text{process\_tweet}(\text{tweet}) \rightarrow [w_1, w_2, w_3, \dots]$
3. Compute  $\text{freq}(w, \text{class})$
4. Get  $P(w \mid \text{pos})$ ,  $P(w \mid \text{neg})$
5. Get  $\lambda(w)$
6. Compute  $\text{logprior} = \log(P(\text{pos}) / P(\text{neg}))$



deeplearning.ai

# Testing Naïve Bayes

# Outline

- Predict using a Näive Bayes Model
- Using your validation set to compute model accuracy



# Predict using Naïve Bayes

- log-likelihood dictionary  $\lambda(w) = \log \frac{P(w|pos)}{P(w|neg)}$

$$\logprior = \log \frac{D_{pos}}{D_{neg}} = 0$$

- Tweet: [I, pass, the, NLP, interview] 🍀

$$score = -0.01 + 0.5 - 0.01 + 0 + \logprior = 0.48$$

$$pred = score > 0$$

word	$\lambda$
I	-0.01
the	-0.01
happi	0.63
because	0.01
pass	0.5
NLP	0
sad	-0.75
not	-0.75

# Testing Naïve Bayes

- $X_{val}$   $Y_{val}$   $\lambda$   $logprior$

$$score = predict(X_{val}, \lambda, logprior)$$

$$pred = score > 0 \quad \begin{bmatrix} 0.5 \\ -1 \\ 1.3 \\ \vdots \\ score_m \end{bmatrix} > 0 = \begin{bmatrix} 0.5 > 0 \\ -1 > 0 \\ 1.3 > 0 \\ \vdots \\ score_m > 0 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 1 \\ \vdots \\ pred_m \end{bmatrix}$$

# Testing Naïve Bayes

- $X_{val}$   $Y_{val}$   $\lambda$   $logprior$

$score = predict(X_{val}, \lambda, logprior)$

$pred = score > 0$

$$\frac{1}{m} \sum_{i=1}^m (pred_i == Y_{val_i})$$

$$\begin{bmatrix} \frac{0}{1} \\ 1 \\ \vdots \\ pred_m \end{bmatrix} == \begin{bmatrix} \frac{0}{0} \\ 1 \\ \vdots \\ Y_{val_m} \end{bmatrix}$$
$$\begin{bmatrix} \frac{1}{0} \\ 1 \\ \vdots \\ pred_m == Y_{val_m} \end{bmatrix}$$

# Summary

- $X_{val} \ Y_{val} \longrightarrow$  Performance on unseen data
- Predict using  $\lambda$  and *logprior* for each new tweet
- Accuracy  $\longrightarrow \frac{1}{m} \sum_{i=1}^m (pred_i == Y_{val_i})$
- What about words that do not appear in  $\lambda(w)$ ?



deeplearning.ai

# Applications of Naïve Bayes

# Applications of Naïve Bayes

$$P(pos|tweet) \approx P(pos)P(tweet|pos)$$

$$P(neg|tweet) \approx P(neg)P(tweet|neg)$$

$$\frac{P(pos|tweet)}{P(neg|tweet)} = \frac{P(pos)}{P(neg)} \prod_{i=1}^m \frac{P(w_i|pos)}{P(w_i|neg)}$$

# Applications of Naïve Bayes

Author identification:

$$\frac{P(\text{Shakespeare}|\text{book})}{P(\text{Shakespeare}|\text{book})}$$

Spam filtering:

$$\frac{P(\text{spam}|\text{email})}{P(\text{nospam}|\text{email})}$$

# Applications of Naïve Bayes

Information retrieval:

$$P(\text{document}_k | \text{query}) \propto \prod_{i=0}^{|\text{query}|} P(\text{query}_i | \text{document}_k)$$

Retrieve document if  $P(\text{document}_k | \text{query}) > \text{threshold}$

"Icon made by [Vector Market](https://www.vectormarket.com) from [www.flaticon.com](https://www.flaticon.com)"



# Applications of Naïve Bayes

Word disambiguation:

$$\frac{P(\text{river}|\text{text})}{P(\text{money}|\text{text})}$$

Bank:



["Pictures with CC"](#)

# Naïve Bayes Applications

- Sentiment analysis
- Author identification
- Information retrieval
- Word disambiguation
- Simple, fast and robust!



deeplearning.ai

# Naïve Bayes Assumptions

# Outline

- Independence
- Relative frequency in corpus

# Naïve Bayes Assumptions

- Independence

“It is sunny and hot in the Sahara desert.”



# Naïve Bayes Assumptions

“It’s always cold and snowy in \_\_\_.”



spring?? summer? fall?? winter??

# Naïve Bayes Assumptions

- Relative frequencies in corpus



# Summary

- Independence: Not true in NLP
- Relative frequency of classes affect the model





deeplearning.ai

# Error Analysis

# Outline

- Removing punctuation and stop words
- Word order
- Adversarial attacks

# Processing as a Source of Errors: Punctuation

**Tweet:** My beloved grandmotherX(

**processed\_tweet:** [belov, grandmoth]

# Processing as a Source of Errors: Removing Words

**Tweet:** This is not good, because your attitude is not even close to being nice.

**processed\_tweet:** [good, attitude, close, nice]

# Processing as a Source of Errors: Word Order

**Tweet:** I am happy because I do not go.



**Tweet:** I am not happy because I did go.



# Adversarial attacks

## Sarcasm, Irony and Euphemisms

**Tweet:** This is a ridiculously powerful movie. The plot was gripping and I cried right through until the ending!

**processed\_tweet:** [ridicul, power, movi, plot, grip, cry, end]

# Summary

- Removing punctuation
- Removing words
- Word order
- Adversarial attacks